Dharama Teja Samudrala

New York, NY 14221
+18573983456 | dharmatejas102@gmail.com | linkedin | github

Summary

Backend Software Engineer with expertise in building scalable, distributed systems and delightful user-facing features. Passionate about designing robust REST APIs and developing maintainable backend solutions using **Golang, Java, and distributed systems technologies**. Proven ability to architect and launch greenfield features from inception to implementation, with a strong focus on performance, reliability, and cross-functional collaboration.

Experience

Glean New York, USA

Backend Software Engineer March 2025 – Present

- Developed and launched scalable backend features for the Work AI platform, architecting **REST APIs** with **Golang** that maximize web client flexibility and support enterprise knowledge graph capabilities.
- Designed distributed system components leveraging microservices architecture, implementing efficient **CRUD functionality** across both SQL and NoSQL databases to support Glean's enterprise search and AI assistant infrastructure.
- Collaborated closely with cross-functional teams including product managers, designers, and data scientists to translate complex requirements into elegant, user-centric technical solutions.
- Implemented robust testing and monitoring frameworks to ensure API reliability, performance, and seamless integration across Glean's enterprise SaaS connectors.

Citi New York, USA

Data Engineer March 2025 – Present

- Engineered scalable ETL pipelines using **Airflow, PySpark, and AWS services**, automating data ingestion from multiple sources and reducing manual processing complexity.
- Developed real-time streaming infrastructure with **Kafka** and **Zookeeper**, processing complex event streams with high reliability and low-latency performance.
- Optimized data storage strategies across cloud platforms, implementing advanced partitioning and compression techniques to enhance query performance and reduce infrastructure costs.
- Implemented CI/CD and monitoring frameworks using **Jenkins** and **AWS CodePipeline**, introducing automated data quality checks that significantly improved pipeline reliability.

Fidelity Investments Raleigh, NC, USA

Data Engineer Jun 2024 – Aug 2024

- Optimized **PySpark** financial processing workflows on **AWS EMR**, implementing advanced DataFrame optimizations that reduced execution time while maintaining data integrity.
- Architected serverless API solutions using **AWS API Gateway** and **Lambda functions**, conducting comprehensive load testing to validate scalability and performance.
- Evaluated and prototyped graph database solutions using **AWS Neptune**, analyzing performance characteristics for complex workflow data models.
- Developed data pipeline automation using **Apache Airflow**, creating sophisticated DAGs with robust task dependencies and error handling.

MetLife India

AWS Data Engineer Dec 2021 – Jul 2023

- Configured and maintained **Amazon S3** data storage infrastructure, implementing advanced partitioning strategies that improved query performance in **AWS Athena**.
- Developed robust **PySpark ETL pipelines** for transforming and enriching transactional data across multiple AWS sources, ensuring high data quality and consistency.
- Implemented event-driven data ingestion using **AWS Lambda** and **SQS**, reducing manual pipeline interventions and increasing automation efficiency.
- Deployed data engineering workflows through automated **CI/CD pipelines** using **Jenkins** and **GitLab Actions**, streamlining deployment processes.

Technical Skills

Languages:	**Golang**, **Java**, Python, SQL
Backend & APIs:	REST APIs, gRPC, OpenAPI, Microservices Architecture
Cloud & Infrastructure:	AWS (S3, Lambda, EMR, API Gateway), Kubernetes, Docker
Databases:	SQL, NoSQL, AWS Neptune, Snowflake, Redshift
Tools & Frameworks:	Apache Airflow, Jenkins, GitLab CI/CD, Kafka, Zookeeper
Data Engineering:	PySpark, ETL, Data Pipelines, Distributed Systems

Projects

Enterprise Work AI Microservices Framework Personal Project

- Designed a scalable microservices communication framework using **Golang**, **gRPC**, and **Kubernetes**, implementing advanced service mesh patterns for enterprise-grade reliability.
- Implemented mutual TLS (mTLS), circuit breakers, and sophisticated retry policies to ensure robust inter-service communication and fault tolerance.
- Built comprehensive observability infrastructure using **Prometheus**, **Grafana**, and **Jaeger**, enabling deep performance insights and distributed tracing.
- Automated deployment with **Docker Compose** and **Helm charts**, demonstrating infrastructure-as-code principles and facilitating seamless scalability.

Distributed Knowledge Graph Prototype Academic Project

- Developed a proof-of-concept knowledge graph system using **Golang** and **AWS Neptune**, exploring advanced graph database techniques for enterprise knowledge management.
- Implemented efficient graph traversal and query optimization strategies, processing complex relationship mappings with sub-millisecond latency.
- Created a modular, extensible architecture supporting multiple data ingestion methods and flexible schema evolution.
- Demonstrated advanced graph database capabilities through practical implementation of entity resolution and semantic relationship detection.

Real-time Event Processing System Personal Project

- Architected a high-performance event streaming platform using **Golang**, **Kafka**, and **Kubernetes**, capable of processing 5,000 events per second with minimal latency.
- Implemented sophisticated event routing and transformation logic, showcasing advanced distributed systems design principles.
- Developed comprehensive monitoring and alerting mechanisms using **Prometheus** and **Grafana** for real-time system observability.
- Achieved 99.99% system reliability through intelligent failure recovery and graceful degradation strategies.

Education

Rochester Institute Of Technology Rochester, NY, USA
Master of Science in Computer Science Aug 2023 – Aug 2025